

STA3105-01 Bayesian Statistics
Homework 5
DUE Friday, November 18

Copying homework solutions from others lead to a 0 score. No late submission is allowed. Your solution should contain both code and a corresponding explanation for the answer. Submit your HW through LearnUs. **You should submit (1) a report file (pdf) and (2) relevant code files.** We will implement MCMC algorithm for a Bayesian logistic regression as follows:

$$\begin{aligned}\mathbf{Y} &\sim \text{Bernoulli}(\mu(\mathbf{X})) \\ \log\left(\frac{\mu(\mathbf{X})}{1 - \mu(\mathbf{X})}\right) &= \mathbf{X}\beta \\ \beta_j &\sim N(0, 10) \text{ for } j = 1, \dots, 4\end{aligned}$$

1. (10 points) Simulate the dataset as follows.
 - (a) Let $\mathbf{X}_i \in \mathbb{R}^4$ be the predictors for i th observation. For $i = 1, \dots, n$, simulate $\mathbf{X}_i \sim N(0, \mathbf{I})$ independently. Here $n = 300,000$ and \mathbf{I} is an identity matrix.
 - (b) For $i = 1, \dots, n$, simulate $Y_i \sim \text{Bernoulli}(\mu(\mathbf{X}_i))$ independently. Set the true regression coefficient value as $\beta = (0.5, -0.5, 0, 1)$.
2. (90 points) Implement the MCMC algorithm for the simulated dataset in Problem 1. Here, you should write down the C++ code using `RcppArmadillo` library as follows.
 - (a) Write down the `Rcpp` function for evaluating joint likelihood function for given $\beta_1, \beta_2, \beta_3, \beta_4, \mathbf{Y}, \mathbf{X}$. Compare the log-likelihood values calculated from your `Rcpp` function with `dbinom` function in R.
 - (b) Write down the `Rcpp` function for evaluating prior for given $\beta_1, \beta_2, \beta_3, \beta_4$.
 - (c) Using the functions defined in (a),(b), write down the `Rcpp` function to generate posterior samples from the above specified model.
 - (d) Report the trace plots, density plots, 95% HPD intervals, posterior mean, acceptance probability, and effective sample size for all parameters. Check whether your MCMC samples can recover the true $\beta = (0.5, -0.5, 0, 1)$ well.
3. (Extra credit: 20 points) Now implement the divide-and-conquer MCMC algorithm to the simulated dataset. Here, you should write down the C++ code using `RcppArmadillo` and `OpenMp` libraries as follows.
 - (a) Divide the entire dataset into S shards; S should be set to the maximum number of cores available on your personal computer. Report the available number of cores using `detectCores` function in `parallel` package.
 - (b) Run S separate MCMC to generate subset posterior samples for each shard. This step should be implemented in parallel via `OpenMp` library.

- (c) Combine subset posterior samples using weighted average as suggested in the consensus Monte Carlo algorithm.

Report the followings:

- (a) Report the user time and elapsed time that can be checked through `proc.time` function in `R`.
- (b) For each subset posterior, report the trace plots, density plots, 95% HPD intervals, posterior mean, acceptance probability, and effective sample size for all parameters.
- (c) Draw the density plot of S subset posteriors and the combined posterior for all parameters. You can overlap all the densities with different colors.